

Laith Tahboub

[linkedin.com/in/laith-tahboub](https://www.linkedin.com/in/laith-tahboub)

ltahboub@umd.edu

+1 (858) 428-5570

EDUCATION

• University of Maryland, College Park

College Park, MD

B.S. in Computer Science, Minor in Mathematics; GPA: 3.61, Dept. Honors, Dean's List

Dec. 2026

Relevant Courses: Machine Learning, Natural Language Processing, Algorithms, Data Science, Programming Languages, Computer Systems, Discrete Structures, Object-Oriented Programming, Linear Algebra, Number Theory

SKILLS

• **Languages:** Python, C, Java, C++, OCaml, SQL, HTML/CSS, JavaScript, L^AT_EX

• **Tools:** PyTorch, GCP, Git, Pandas, NumPy, LangGraph, Google ADK, PostgreSQL, Django, AWS, Spark

EXPERIENCE

• Incorta

SF Bay Area, CA

July 2025 – Oct. 2025

Machine Learning Engineer Intern

- Created a new sub-agent for the company's AI data assistant (Incorta Copilot) using Google ADK and Agentspace. 50,000+ employees from Broadcom, Starbucks, Comcast, and more use this AI agent daily.
- Architected an automated benchmarking framework for Incorta Copilot, now used by 30+ salespeople in 1,000+ customer demos yearly. Tool training integrated into company-wide sales curriculum.
- Engineered a high-accuracy employee resignation prediction model for the Saudi Stock Exchange (12th largest globally) using gradient boosting, achieving 0.94 ROC-AUC and 0.90 F1-score.
- Developed explainable anomaly detection methods; presented to 50+ engineers and executives. Recognized by leadership and integrated into Copilot, used by Fortune 500 companies managing billions in revenue.

• University of Maryland

College Park, MD

Teaching Assistant — CMSC 250 Discrete Structures

Sept. 2025 – Present

- Lead weekly discussion sessions for 32 students covering topics such as mathematical proof techniques, logic, set theory, graph theory, and combinatorics, resulting in improved student comprehension and scores.
- Grade assignments and exams for 220+ students, providing detailed feedback on mathematical reasoning and proof writing while maintaining consistent evaluation standards across multiple TA sections.
- Hold regular office hours to provide individualized support for challenging concepts (e.g. induction, inference, number theory) helping students develop mathematical maturity for advanced CS coursework.

Research Assistant — Professor Mohammed Nayeem Teli's NLP Lab

Dec. 2024 – June 2025

- Led a cross-functional team of 5 undergraduate researchers in building an evaluation pipeline for LLM alignment to normative text corpora, coordinating experiments and managing HPC workflows.
- Created a novel evaluation method generating 100K+ QA pairs from philosophical, legal, and religious texts to benchmark top LLMs (GPT-4, Claude, Llama 3) using PyTorch, HuggingFace, PEFT, and SLURM.
- Discovered that top LLMs are consistently more factually aligned than morally aligned, with results forming the basis of a planned NAACL 2026 publication.

• U.S. News and World Report

Remote Contract via UMD App Dev

Sept 2024 – Jan 2025

Data and Backend Engineer

- Selected from a competitive pool of candidates to develop and deploy an internal tool that visualizes and organizes API data, supporting the company's quality assurance workflows.
- Collaborated directly with U.S. News engineers to create the backend using FastAPI and integrated Scikit-learn anomaly detection with a KNN model, ensuring scalable performance and robust data analysis.
- Delivered a production-ready platform now used daily by the U.S. News QA team to validate datasets, improving anomaly detection efficiency and reducing manual review time.

PROJECTS

• Stilltheory.com

May 2022 – Present

- Launched StillTheory, a web app for chess opening analysis now used by 500+ registered players, providing personalized repertoire insights and real-time visual feedback.
- Built a responsive backend infrastructure using Celery, Django Channels, Redis, and WebSockets to enable live progress updates, supporting concurrent usage across 40+ countries without performance degradation.
- Engineered a data pipeline that ingests and processes millions of chess game records from external APIs, transforming them into personalized opening trees per user in under 30 seconds.

• Predicting Professor Ratings

Feb. 2025

- Built a machine learning model to predict professor ratings on PlanetTerp using features such as class difficulty, department, number of courses taught, tenure, and review sentiment.
- Trained Random Forest, Multinomial Logistic Regression, and sentiment-only models, finding that average review sentiment alone enabled 88% accuracy, and that mid-range ratings were the hardest to classify.